ETHICAL DILEMMA

CRISIS! Your data scientist discovered hidden bias in the AI.

THE SITUATION: It's Week 6 of the pilot. Everything looks good: - AI accuracy: 91% - Customer satisfaction: 82% - Team is engaged - CEO is happy

THEN: Your data scientist comes to you with a problem:

> *"I was analyzing our training data. We trained the AI on 5 years of historical customer service records. That data is biased.*
>
> *Older customers get faster response times. Rural customers get slower. Specific zip codes marked as 'high-risk' get deprioritized.*
>
> *The AI learned this bias from the patterns. It's now recommending: Give priority to urban customers, flag rural customers as 'high-effort,' speed responses to wealthy zip codes.*
>
> *We can fix this by using a different algorithm, but accuracy drops to 85%. OR we ship it as-is—customers won't notice, metrics stay good, and the bias is mild."*

YOUR DECISION POINT

Do you fix the bias (accuracy drops to 85%) or ship the biased system (keep 91% accuracy)?

YOUR TASK (30 minutes to decide)

1. What's the real trade-off here? (It's not what it looks like)

2. What are your options? (List at least 3)

3. What would you recommend and why?

4. How do you communicate this decision to the CEO?

5. What changes in how you design AI projects after this?