PROJECT SCOPING FRAMEWORK FOR AI PILOTS

AI Leadership & Project Management Masterclass

Curtin University

2025-10-29

PROJECT SCOPING FRAMEWORK FOR AI PILOTS

The Goldilocks Principle: Not Too Big, Not Too Small, Just Right

---

When to Use This Framework

Use this framework when:

- **Starting an AI pilot** - define what you're actually building
- **Scope creep emerges** - pushback on "just one more feature"
- **Timeline questions arise** - how do we know when we're done?
- **Resource constraints hit** - what can we cut without breaking the project?
- **Stakeholder disagreements** - different people want different things

---

The Three Core Principles

1. THE GOLDILOCKS PRINCIPLE: NOT TOO BIG, NOT TOO SMALL

Too Small: Scope so narrow that even success doesn't prove anything - Example: "AI chatbot for 5 questions, one team, 2 weeks" - Problem: If it succeeds, you learn nothing valuable. If it fails, was it the AI or the tiny scope?

Too Big: Scope so ambitious that you're guaranteed to fail or take years - Example: "AI for all customer service, all query types, all 50 locations" - Problem: Too many unknowns, too much risk, impossible to learn incrementally

Just Right: Scope that tests the key assumptions with acceptable risk - Example: "AI chatbot for 30% of queries (order tracking + returns) at one location for 6 weeks" - Problem solved: Proves AI can work in customer service, teaches you about real-world data quality, manageable risk

---

2. SUCCESS METRICS: DEFINE "DONE" BEFORE YOU START

Don't start until you can answer:

Business Metrics

- What does success look like financially?
- What's the ROI you're targeting?
- Cost per query? Revenue uplift? Time saved?

Example targets: - Reduce response time from 26 hours to <4 hours - Reduce cost per query from $18 to $12 - Handle 30% of queries without human intervention

Quality Metrics

- How accurate does the AI need to be?
- What's your acceptable error rate?
- What kind of errors are acceptable vs. unacceptable?

Example targets: - Accuracy 90% on routine queries - <2% escalation rate (queries that need human review) - 80% customer satisfaction with AI responses

Adoption Metrics

- Who needs to use this system?
- How will you know they're actually using it?
- What's your adoption target?

Example targets: - 80% of eligible queries routed to AI - Team morale doesn't decline - Zero voluntary opt-outs from using the system

Timeline Metrics

- When do you evaluate?
- How long does the pilot run?
- How many data points before you decide?

Example targets: - 6-week pilot (long enough to find real issues, short enough for momentum) - Evaluate weekly during pilot - Decision point: End of week 6

---

3. RISK MITIGATION: KNOW YOUR RED LINES

Identify risks upfront:

Technical Risks

- "What if data quality is worse than expected?" → Plan for data cleaning budget
- "What if accuracy is lower than needed?" → Plan for manual QA phase
- "What if the model doesn't work?" → Have a fallback plan (human handling)

People Risks

- "What if the team resists using AI?" → Plan for change management + communication
- "What if we lose key staff?" → Cross-train, document everything
- "What if leadership loses patience?" → Weekly updates with clear metrics

Ethical Risks

- "What if the AI discriminates?" → Plan for bias testing before scale
- "What if customer data is exposed?" → Security review before launch
- "What if we can't explain decisions?" → Interpretability checks built in

Business Risks

- "What if it costs more than we budgeted?" → Reserve 20% contingency
- "What if it takes longer than planned?" → Know your kill criteria
- "What if business priorities shift?" → Define what happens then

---

The Scoping Checklist

Before you start your AI pilot, confirm:

Scope Definition

☐ **One clear problem** - not five problems solved with one AI
☐ **One clear user group** - not "everyone in the company"
☐ **Clear boundaries** - what the AI will handle vs. what it won't
☐ **Realistic timeline** - 6-12 weeks is typical for pilots (not 2 weeks, not 2 years)

Success Criteria

☐ **Business metrics defined** - what does success look like financially?
☐ **Quality metrics defined** - accuracy, error rates, customer satisfaction
☐ **Adoption metrics defined** - how will people use this?
☐ **Go/No-Go criteria documented** - when do we scale vs. pivot vs. kill?

Risk Management

- ☐ **5+ risks identified** - technical, people, ethical, business
- ☐ **Mitigation plan for each** - not "hope it doesn't happen"
- ☐ **Red lines defined** - what would make us kill the project?
- ☐ **Contingency budget allocated** - 15-20% of total budget

Stakeholder Alignment

- ☐ **CEO agrees on timeline and ROI expectations**
- ☐ **CFO approves budget and contingency**
- ☐ **Data team agrees timeline is realistic**
- ☐ **Operations team agrees on adoption approach**
- ☐ **Security team approves data handling**

Common Scoping Mistakes to Avoid

Mistake 1: "Let's just see what happens"

Problem: No success criteria = everything is success **Fix:** Define at least 3 metrics before you start

Mistake 2: "We'll figure out the risks as we go"

Problem: Risks discovered mid-project derail timelines **Fix:** Identify and mitigate risks upfront

Mistake 3: "We want the MVP but also X and Y and Z"

Problem: Scope creeps until it's not an MVP anymore **Fix:** Say "no" to features. You can do them in phase 2.

Mistake 4: "We'll decide if it's working when it's working"

Problem: No clear decision point = project limbo **Fix:** Set a calendar date: "Week 6, we decide: Scale, Pivot, or Kill"

Mistake 5: "Success means the AI works technically"

Problem: AI might work but team won't use it, CEO won't fund it, ethics team will block it **Fix:** Success = technical + adoption + business + ethical

---

How to Use This in Your Projects

Week 1: Define scope, success metrics, risks, go/no-go criteria (you did this!)

Weeks 2-6: Run the pilot, collect data, watch metrics

Week 6: Gather everyone, look at data, decide: **Scale, Pivot, or Kill**

- **Scale:** Metrics hit targets, no showstopper risks, team is ready → Expand to next phase
- **Pivot:** Some metrics hit, some don't, but clear path forward → Adjust and extend pilot
- **Kill:** Fundamental assumptions wrong, unresolvable risks → Stop and try something else

Key insight: "Kill" is not failure. It's data. You learned something valuable by spending $150K on a 6-week pilot instead of $500K on a failed year-long rollout.

Remember

The best scope is: - **Small enough** to learn quickly - **Big enough** to prove something valuable - **Clear enough** to know when you're done - **Flexible enough** to adapt when reality surprises you

Your job isn't to predict the future perfectly. Your job is to learn the fastest way possible so you can make good decisions.