

SCALE OR KILL EXERCISE

AI Leadership & Project Management Masterclass

Curtin University

2025-10-27

SCALE OR KILL EXERCISE

Three AI Project Case Studies

INSTRUCTIONS:

Your group will analyze three AI pilot projects and decide whether to SCALE, PIVOT, or KILL each one. Be prepared to defend your decisions.

CASE STUDY A: THE CLEAR SUCCESS

AI Inventory Prediction System - Fashion Retailer

PROJECT OVERVIEW

Company: StyleMart - Mid-sized fashion retailer (35 stores)

The Problem:

StyleMart was losing money on both ends of inventory management: - Stockouts (lost sales when popular items sold out) - Overstock (heavy markdowns on items that didn't sell)

The Solution:

AI system that predicts demand for clothing items by store, using: - Historical sales data - Weather patterns - Local events (concerts, sports, festivals) - Social media trends - Competitor pricing

The system automatically adjusts stock orders for each store.

Timeline: 12-month pilot across 10 of 35 stores

INVESTMENT

Category	Amount
AI Platform & Development	\$180,000
Implementation & Integration	\$45,000
Staff Training	\$15,000
Total Investment	\$240,000

RESULTS AFTER 12 MONTHS

Inventory Performance Metrics:

Metric	Before AI	After AI	Change
Stock-out rate	45%	12%	↓ 73%
Overstock rate	31%	9%	↓ 71%
Inventory turnover ratio	4.2x per year	6.8x per year	↑ 62%
Waste from unsold seasonal items	\$420K	\$176K	↓ 58%

Financial Impact:

Impact Area	Annual Value
Revenue increase (fewer stockouts)	+\$1,200,000
Cost savings (reduced overstock)	+\$800,000
Total Annual Benefit	\$2,000,000
ROI (First Year)	733%


Operational Improvements:

- **Ordering process time:** Reduced by 70% (from 8 hours/week to 2.5 hours/week per store manager)
- **Merchandising team:** Freed up 15 hours/week to focus on strategy vs. manual forecasting
- **Planning accuracy:** Improved from 62% to 87%

STAKEHOLDER FEEDBACK

Store Managers (8 of 10 positive): > “Finally! We get the right products at the right time. I’m not constantly calling head office asking for stock transfers.” - Manager, Bondi Store

“First month was rough - the AI made some weird predictions. But after it learned our store’s patterns, it’s been amazing.” - Manager, Perth Store

Two skeptics initially: > “I didn’t trust it for the first 3 months. Kept overriding its  Curtin University recommendations. Then I realized my gut instinct was wrong more often than the AI.” - Manager, Melbourne Store (now biggest advocate)

Merchandising Team: > “This has transformed how we work. We’re now analyzing trends and planning strategy instead of drowning in spreadsheets.” - Head of Merchandising

CFO: > “The ROI is undeniable. This is exactly the kind of AI investment that makes sense.”

Customers: - Customer satisfaction: +14% (products in stock when they want them) - Net Promoter Score: +8 points - Fewer complaints about “never having my size”

TECHNICAL PERFORMANCE

- **Prediction accuracy:** 87% (vs. 62% with old manual forecasting)
 - **System uptime:** 99.7%
 - **Integration:** Seamless with existing ERP and POS systems
 - **Maintenance:** 2 hours/month (mostly automated)
-

CHALLENGES ENCOUNTERED

Month 1-2: - Accuracy was only 68% (AI was learning store-specific patterns) - Store managers didn’t trust it, kept overriding recommendations - One store had a stock disaster (AI didn’t account for local road closure during festival)

Month 3: - One store manager resisted for 3 months (“I know my customers better than a computer”) - Eventually came around after seeing results at other stores

Month 7: - IT needed to upgrade server capacity - unexpected \$12,000 cost

By Month 9: - All issues resolved, system performing consistently

CURRENT STATUS

- **All 10 pilot stores** enthusiastically want to keep it
- **15 other stores** have formally requested access
- **Merchandising VP** champions it in every executive meeting
- **Store managers** sharing success stories internally

Team Recommendation: Scale to all 35 stores immediately

YOUR ANALYSIS

Decision: SCALE PIVOT KILL

If SCALE:

Rollout timeline:

Additional investment needed:

Key risks to manage during scaling:

- 1.
- 2.
- 3.

Success criteria for full rollout:

What could go wrong at scale that didn't go wrong in pilot?

If PIVOT:

What would you change?

Why pivot instead of scale or kill?

If KILL:

Why would you kill a project with 733% ROI?

What are you seeing that others aren't?

Questions You Still Have:

- 1.
 - 2.
 - 3.
-

CASE STUDY B: THE CLEAR FAILURE

AI Staff Scheduling System - Restaurant Chain

PROJECT OVERVIEW

Company: CasualBite - Fast-casual restaurant chain (24 locations)

The Problem: - Labor costs were 28% of revenue (target: 26%) - Inefficient scheduling (overstaffed at slow times, understaffed during rushes) - Manual scheduling took managers 6 hours/week

The Solution: AI creates optimal staff schedules based on predicted customer traffic, considering: - Historical traffic patterns - Weather forecasts - Local events - Holidays and school calendars - Staff availability and preferences

Goal: Reduce labor costs while maintaining (or improving) service quality

Timeline: 6-month pilot across 8 restaurants

INVESTMENT

Category	Amount
Software License (Annual)	\$95,000
Implementation & Integration	\$60,000
Training	\$25,000
Total Investment	\$180,000

RESULTS AFTER 6 MONTHS

Labor Metrics:

Metric	Before AI	After AI	Change
Labor cost as % of revenue	28%	26%	↓ 2%
Total staff hours per week (all 8 locations)	2,840 hours	2,520 hours	↓ 11%
Achieved cost savings goal!			Met target

Looks successful, right? But...

Service Impact:

Metric	Before AI	After AI	Change
Average customer wait time	8 minutes	18 minutes	↑ 125%
Customer complaints (per month)	23	78	↑ 340%
Table turnover rate	2.8 turns/table/dinner	2.2 turns/table/dinner	↓ 22%
Customer satisfaction score	82%	61%	↓ 21 points
Walkouts (customers who left before being served)	~2 per week	~8 per week	↑ 400%

Staff Impact:

Metric	Before AI	After AI	Change
Employee turnover (annual rate)	23%	67%	↑ 191%
Staff satisfaction survey	68%	31%	↓ 37 points
Shifts scheduled with <24hrs notice	12/month	54/month	↑ 450%
“Split shifts” assigned	8/month	30/month	↑ 275%
Staff calling in sick (suspected protest)	4% of shifts	11% of shifts	↑ 175%

WHAT WENT WRONG

The AI optimized for the wrong objective.

It successfully minimized labor costs, but in doing so, it created:

1. Unpredictable Schedules

- Staff received schedules with <24 hours notice
- Shifts changed frequently based on AI weather predictions
- Impossible for staff to plan childcare, second jobs, or personal lives
- Example: “You’re scheduled Thursday 11am-2pm... wait, now you’re off Thursday and working Friday instead.”

2. Chronic Understaffing

- AI predicted lower traffic than actually occurred
- Scheduled minimum staff “just in case” traffic was low
- When restaurants got busy, staff were overwhelmed
- No buffer for unexpected rushes

3. Split Shifts

- AI assigned shifts like: Work 11am-2pm (lunch), then come back 6pm-9pm (dinner)
- Saves labor cost (don’t pay for 2-6pm gap)
- Staff hated having to come in twice with unpaid gap in between

4. Zero Consideration for Employee Wellbeing

- Algorithm saw staff as interchangeable units
- No consideration for preferences, fairness, or morale
- Optimized for cost, ignored human factors

REAL EXAMPLES

Example 1: The Saturday Disaster - AI predicted low traffic for Saturday dinner (based on weather forecast of rain) - Scheduled only 3 servers and 2 kitchen staff - Weather cleared up, restaurant got slammed - 45-minute wait times, terrible service, 12 groups walked out - Saved \$240 in labor that night, lost ~\$1,800 in revenue - Two servers quit the next week from stress

Example 2: The Single Mom - Server with consistent childcare needs (needed 9am-5pm shifts) - AI started scheduling her for split shifts and evening close - She couldn’t find childcare, had to quit - Lost a 4-year employee with excellent reviews

Example 3: The Text Message - Staff member received text at 8pm: “Your shift tomorrow has been cancelled due to predicted low traffic” - Had already arranged childcare and turned down other work - Lost a day’s wages with 14 hours notice

FINANCIAL REALITY CHECK

What the CFO saw initially: - Labor cost savings: \$84,000 over 6 months

What the full picture shows: - Labor savings: +\$84,000 - Revenue loss (slower table turns, walkouts, bad reviews): -\$127,000 - Recruiting & training costs (67% turnover): -\$93,000 - Brand damage (one-star reviews): -\$45,000 (estimated) - **Net Financial Impact: -\$181,000 LOSS**

STAKEHOLDER FEEDBACK

Restaurant Managers (7 of 8 want it gone): > “This system is destroying my team. I’m losing good people who’ve been with us for years because they can’t plan their lives.” - Manager, Brisbane Location

“We’re saving money on labor but hemorrhaging revenue because service is terrible. This is penny-wise, pound-foolish.” - Manager, Sydney Location

“I’ve gone back to manual scheduling. I’d rather spend 6 hours a week on schedules than deal with this mess.” - Manager, Melbourne Location (already stopped using the AI)

Staff (via anonymous survey): > “I have two jobs. I can’t work when the AI changes my schedule with 12 hours notice. I’m looking for a new job.”

“I used to love working here. Now I dread checking my phone to see what crazy schedule the computer made.”

“The AI treats us like robots. Management clearly doesn’t care about us anymore.”


District Manager: > “We’re saving on labor costs but losing customers. We’ve dropped from 4.2 stars to 2.8 stars on Google. Not worth it.”

HR Director: > “We’re hemorrhaging talent. Exit interviews all mention the scheduling system. It’s going to cost us more to rebuild the team than we saved.”

Customers (Online Reviews): > “Used to be my favorite restaurant. Now there’s never enough staff. 30-minute wait for fast-casual food? No thanks.”

“You can tell the staff are stressed and understaffed. Not their fault, but I won’t be back.”

TECHNICAL PERFORMANCE



Ironically, the AI worked perfectly: - Algorithm functioned as designed - Minimized labor costs - System uptime: 98% - Integration with scheduling software: Flawless - Predictions were actually quite accurate

The problem: It was optimized for the wrong goal.

The AI did exactly what it was told to do (minimize labor cost). It just ignored all the human and business consequences.

CURRENT STATUS

- **6 of 8 restaurants** have formally requested to stop using it
- **2 managers** already reverted to manual scheduling (without approval)
- **Multiple formal HR complaints** filed
- **3 threatened lawsuits** from employees (unfair scheduling practices)
- **Union is getting involved** (in locations where staff are unionized)
- **Customer satisfaction declining** at all pilot locations
- **Google/Yelp ratings down** an average of 1.4 stars

Team Recommendation: Kill the project immediately

YOUR ANALYSIS

Decision: SCALE PIVOT KILL


If KILL:

What's the key lesson learned?

What would you do differently if you tried this again?

How do you communicate this to stakeholders?

If PIVOT:



What would you change about the approach?

Would changing the optimization criteria fix this?

For example, what if AI optimized for: - Minimize labor cost AND maintain customer satisfaction >80%
- Minimize labor cost AND maintain staff satisfaction >70% - Minimize labor cost AND limit schedule changes to 72hrs notice

Would that address the core issues?

Is this worth pivoting or is the concept fundamentally flawed?

If SCALE:

Why would you scale this?

What are you seeing that the restaurants aren't?

Reflection Questions:

What does this case teach about: - Choosing the right success metrics? - AI ethics and human impact? - The difference between “working as designed” and “working well”?

CASE STUDY C: THE AMBIGUOUS ONE

AI Dynamic Pricing System - E-Commerce Platform

PROJECT OVERVIEW



Company: HomeGoods Online - E-commerce retailer (home goods, furniture, decor)

The Problem: - Competitive pricing pressure - Inventory sitting too long (carrying costs) - Revenue optimization opportunities being missed - Manual pricing decisions slow and inconsistent

The Solution: AI adjusts product prices in real-time based on: - Demand signals (page views, cart additions, search volume) - Competitor pricing (monitoring 12 major competitors) - Inventory levels (reduce prices on slow-moving stock) - Time of day, day of week patterns - Customer browsing behavior (how long they looked, how many times they returned)

Prices can fluctuate throughout the day, similar to airline pricing.

Timeline: 9-month pilot on 500 products (out of 12,000 total catalog)

INVESTMENT

Category	Amount
AI Platform Development	\$220,000
Implementation & Integration	\$45,000
A/B Testing Infrastructure	\$35,000
Total Investment	\$300,000

RESULTS AFTER 9 MONTHS

Revenue Metrics:

Metric	Control Group (No AI)	AI Pricing Group	Difference
Revenue per product	Baseline	+8.2%	Significant increase
Gross margin	Baseline	+3.1%	Better margins
Conversion rate	12.8%	12.5%	↓ 2.4% Slightly down
Average order value	\$147	\$164	↑ 11.7% Higher
Units sold per product	Baseline	+2.1%	Slightly up

Financial Impact:



- **Additional revenue:** +\$890,000 over 9 months
- **Annualized projection:** ~\$1,180,000 additional revenue
- **Additional gross margin:** +\$312,000 over 9 months
- **Annualized margin:** ~\$415,000
- **ROI (first year):** 138% (if sustained)

From a pure financial perspective: **This is working.**

BUT... CUSTOMER EXPERIENCE METRICS:

Metric	Before AI	After AI	Change
Customer complaints about pricing	8/month	27/month	↑ 340%
Negative social media mentions	12/month	62/month	↑ 520%
Cart abandonment rate	58%	71%	↑ 13 points
Customer satisfaction score	79%	67%	↓ 12 points
Return rate	12%	18%	↑ 50%
“Price complaint” customer service tickets	45/month	178/month	↑ 396%

From a customer experience perspective: **This is causing problems.**

WHAT'S HAPPENING - REAL SCENARIOS

Scenario 1: The “Price Watcher”

Monday 10am: - Customer browses decorative lamp, price is \$89.99 - Adds to cart but doesn't purchase
- Leaves website

Monday 2pm: - Customer returns to complete purchase - Price is now \$107.99 (+20%) - AI detected high interest, raised price - Customer feels manipulated, abandons cart - Posts on social media: “They're tracking me and punishing me for looking!”

Result: Lost sale + bad publicity

Scenario 2: The “Bad Luck Shopper”



Curtin University

Customer A - Monday morning: - Buys throw pillows: \$79.99 - Happy with purchase

Customer B (Customer A’s friend) - Tuesday afternoon: - Buys identical throw pillows: \$64.99 - Posts on Instagram about the “great deal”

Customer A sees friend’s post: - Realizes they paid \$15 more for identical product 24 hours earlier - Feels ripped off, demands price match - Returns item, writes negative review - Posts complaint: “They charge different prices to different people!”

Result: Lost customer loyalty + negative word of mouth

Scenario 3: The “Social Media Firestorm”

Customer tweets: > “I’ve been watching this coffee table for a week. Price changes EVERY TIME I look at it. \$245 → \$289 → \$227 → \$312. Are they tracking me?? This feels predatory. @HomeGoodsOnline”

Tweet goes viral (12,000 retweets)

Media picks it up: > “HomeGoods Online Uses AI to Track Shoppers and Manipulate Prices”

Reality: AI adjusts by time/demand, NOT by individual customer. But perception = reality.

Scenario 4: The “Algorithm Winner”

Thursday 2am: - Item usually \$150 - AI detects very low demand overnight - Reduces price to \$112 to test price elasticity

Customer finds it: - Thrilled with deal, buys 3 units - Tells 5 friends about amazing price

Friends check 6 hours later: - Price back to \$145 - All 5 friends feel deceived - Brand looks inconsistent and untrustworthy

CUSTOMER FEEDBACK (QUALITATIVE)

Positive Comments (15% of feedback):

“I got an amazing deal on exactly what I needed!”

“Prices seem to drop when items aren’t selling well - good for bargain hunters”

Negative Comments (85% of feedback):



Curtin University

Theme 1: Feeling Tracked (43% of complaints) > “It’s creepy that prices change when I look at products”

“This is predatory pricing - they’re tracking my behavior”

“I feel manipulated, like they know I really want it so they raise the price”

Theme 2: Inconsistency/Unfairness (31% of complaints) > “Prices change every time I look - I can’t trust them”

“My friend paid less for the same thing - that’s not fair”

“How can the price go up \$40 in 3 hours?”

Theme 3: Trust Erosion (28% of complaints) > “I used to love this store but now I feel like they’re trying to scam me”

“Can’t trust a company that changes prices randomly”

“This has destroyed my loyalty to this brand”

MEDIA ATTENTION

Local news story: > “Perth-based HomeGoods Online Uses AI to Charge Different Customers Different Prices”

The article was technically inaccurate (AI adjusts by time/demand, not by individual customer) but:

- Damage to brand perception was real - 47,000 views on the news website - Shared widely on social media
- Comments section filled with anger

Company response was defensive: > “Our dynamic pricing is similar to airlines and ride-sharing services”

Public reaction: > “Yeah, and we hate that too!”

INTERNAL STAKEHOLDER VIEWS



CFO (Pro-Scaling): > “Revenue is up 8.2%. Margin is up 3.1%. ROI is 138%. This is working. We need to scale this to all 12,000 products immediately. Customer complaints are just noise - people always complain. Look at the data.”

CMO (Anti-Scaling): > “Customer satisfaction is down 12 points. That’s catastrophic for long-term brand value. We’re destroying trust for short-term revenue gains. Customer lifetime value is more important than this quarter’s revenue. This is incredibly short-sighted.”

Head of E-commerce (Conflicted): > “The data shows it’s working from a revenue perspective. But the customer feedback is genuinely concerning. I worry we’re winning battles and losing the war. Maybe we’re just not communicating it well?”

Customer Service Director (Strongly Against): > “My team is drowning in pricing complaints. This has created a massive operational burden. Customers are angry, confused, and feel cheated. Whatever revenue we’re gaining, we’re paying for it in customer service costs and brand damage.”

Data Scientist (Neutral/Technical): > “The algorithm is performing exactly as designed. It’s optimizing for revenue. If the business wants to optimize for something else - like customer satisfaction or price stability - we can adjust the model. But right now it’s doing what it was asked to do.”

CEO (Waiting for Recommendation): > “I need a clear recommendation from this group. Is this good for RetailFlow long-term? I don’t care about this quarter - I care about the next 3 years. Make a recommendation you can defend.”

TECHNICAL PERFORMANCE

- **Algorithm works perfectly:** Sophisticated, well-designed
- **Revenue optimization:** Achieving its goal
- **System uptime:** 99.2%
- **Pricing adjustments:** Smooth, accurate, responsive
- **No technical failures**

The question isn’t whether it works technically. The question is whether it’s the right thing to do.

MARKET CONTEXT

What competitors are doing:

- **Amazon:** Uses dynamic pricing extensively (hundreds of price changes per day on some items)
- **Airlines:** Dynamic pricing standard practice
- **Uber/Lyft:** Surge pricing well-established
- **Many retailers:** Moving toward dynamic pricing

BUT:

Some competitors are going the opposite direction: - **Everlane:** Markets “transparent pricing” and “no sales” as competitive advantage - **Costco:** Stable pricing as brand promise - **Some boutique brands:** “Honest pricing, no games” positioning

Regulatory environment: - Australian Consumer Law: No explicit prohibition on dynamic pricing - Growing scrutiny: Politicians discussing regulations - EU: Considering transparency requirements - Consumer advocacy groups: Calling dynamic pricing “exploitative”

THE DILEMMA

Financially successful: 8.2% revenue increase, positive ROI, meeting business objectives

Customer experience suffering: 12-point satisfaction drop, trust erosion, negative sentiment

Long-term impact unclear: - Will revenue gains sustain if brand trust continues to erode? - Are we winning revenue today but losing loyalty tomorrow? - Is 8% revenue gain worth 12-point satisfaction drop? - What happens when competitors do the same thing?

CURRENT STATUS

- **Finance wants to SCALE:** “The numbers are great”
- **Marketing wants to KILL:** “We’re destroying the brand”
- **E-commerce is CONFLICTED:** “It’s working but I’m worried”
- **Customer service wants to KILL:** “This is creating chaos”
- **CEO wants YOUR recommendation**

Team Recommendation: ???

YOUR ANALYSIS (This is the hard one)

Decision: SCALE PIVOT KILL

RISK ASSESSMENT

Consider each decision path:

If You SCALE	If You PIVOT	If You KILL
--------------	--------------	-------------



Best Case:

Worst Case:

Most Likely:

YOUR RECOMMENDATION

We recommend:

Our rationale:

What we're prioritizing:

What we're willing to sacrifice:

IF SCALE:


How will you address customer satisfaction concerns?

What guardrails will you put in place?

How will you communicate this to customers?

Success metrics for the scaled version:

IF PIVOT:



What specific changes would you make?

Examples to consider: - Limit price fluctuation to $\pm 15\%$? - Only adjust prices once daily, not throughout day? - Show customers historical pricing? - Grandfather previous prices for returning customers? - Opt-in "price watch" notifications?

Your pivot approach:

Would this pivot address the core issues?

IF KILL:

How do you justify killing a project with positive ROI?

What's more important than short-term revenue?

How do you respond to the CFO who wants to scale?

DEFENDING YOUR DECISION

How will you defend this to:

The CFO (who sees 8.2% revenue increase):

The CMO (who sees 12-point satisfaction drop):

The CEO (who wants long-term brand value):

Customers (if they ask about your pricing):



THE HARD QUESTIONS

Is dynamic pricing inherently unethical? - Amazon does it - Airlines do it - Is it different when AI does it vs. humans?

What's the difference between "smart pricing" and "price manipulation"?

How do you balance: - Shareholder value (maximize revenue) - Customer trust (stable, fair pricing) - Competitive necessity (others are doing this)

If you kill this, are you leaving money on the table?

If you scale this, are you sacrificing long-term brand for short-term gains?

REFLECTION

What makes this case different from Cases A and B?

What decision framework did you use?

How did your group handle disagreement?

Are you comfortable defending your decision publicly?

FACILITATOR DEBRIEF POINTS



This case deliberately has no clear “right” answer.

Strong responses will: - Acknowledge the tension between financial and experiential metrics - Consider long-term vs. short-term impacts - Think about brand positioning and values - Consider competitive context - Make a clear decision with sound reasoning

Watch for: - Groups that only look at financial metrics - Groups that ignore the positive ROI - Groups that can’t make a decision - Groups that find creative middle ground

Key discussion: - How many groups chose differently? - What values/priorities drove different decisions? - Is there a “right” answer? - What would change your mind?